

RULES

- **No books, no notes, and no calculators.**
- **No bathroom breaks** until after you have completed and submitted the exam.
- **All phones must be completely silent** for the duration of the exam, so please *turn off your phone now!*

Cheating will not be tolerated. If there is any indication that a student may have given or received unauthorized aid on this test, the case will be referred to the Office of the Dean of Students. When you finish the exam, you must sign the following pledge:

“On my honor as a student I, _____, have neither given nor received unauthorized aid on this exam.” (print name clearly)

Signature: _____ Date: November 11, 2022

Page:	2	3	4	5	6	7	8	Total
Points:	12	10	15	21	12	18	12	100
Score:								

Instructions. Answer the following multiple choice questions by selecting all correct choices. If a question has more than one correct choice, it will say in parentheses how many items you should select. *Select all correct choices to receive full credit!*

1. (6 pts) **Big data properties.**

- (a) “Big Data” concerns which of the following types of data?
 - ☐ structured ☐ semi-structured ☐ unstructured ☐ all of these
- (b) JSON and XML are examples of which type of data?
 - ☐ structured ☐ unstructured ☐ semi-structured ☐ none of these
- (c) Which of the following statements is true of unstructured data?
 - ☐ It is generally easier to analyze than other types of data.
 - ☐ It fits neatly into a schema.
 - ☐ It is the most widespread type of data.
 - ☐ It is usually found in tables.

2. (6 pts) **Hardware and Architecture.**

- (a) What kind of hardware is typically used for big data applications?
 - ☐ high-performance supercomputers
 - ☐ low-cost, commodity hardware
 - ☐ dumb terminals
 - ☐ floppy disks
- (b) What is “commodity” hardware?
 - ☐ high-performance supercomputers
 - ☐ discarded, second-hand, or recycled hardware
 - ☐ hardware used for trading commodities (e.g., gold, silver, soy-beans)
 - ☐ generic, low-specification, industry-grade hardware
- (c) Which of the following is **not** a drawback of traditional relational database management system (or RDBMS) when used for big data applications?
 - ☐ They do not make it easy to handle massive volumes of unstructured or semi-structured data.
 - ☐ They require more processors and memory to scale up to big data applications.
 - ☐ They are relatively slow when used to perform SQL queries on large structured data tables.
 - ☐ They do not make it easy to capture and process unstructured or semi-structured data arriving at high velocity.

3. (10 pts) **Programming Paradigms**

(a) Which of the following are programming paradigms? (select three)

- ☐ currying ☐ declarative ☐ dysfunctional ☐ functional ☐ object-oriented

(b) Which of the following characteristics are typical of imperative programs. (select two)

- ☐ variables are *mutable*; their values may change or “mutate”
☐ for loops are usually preferred in favor of recursive function calls
☐ functions are *referentially transparent*
☐ functions are “pure” (do not have *side-effects*)

(c) Which of the following characteristics are typical of functional programs. (select two)

- ☐ variables are *mutable*; their values may change or “mutate”
☐ for loops are usually preferred in favor of recursive function calls
☐ functions are *referentially transparent*
☐ functions are “pure” (do not have *side-effects*)

(d) By definition, a *higher-order function* is a function which

- ☐ is passed as an argument to other functions.
☐ is returned as output by other functions.
☐ is called a higher-order number of times in comparison to “lower-order” functions.
☐ requires a higher-order amount of time to compute in comparison to “lower-order” functions.
☐ accepts a function (or functions) as input or returns a function (or functions) as output.

(e) An expression **e** is called *referentially transparent* provided

- ☐ when reduced to “normal form” **e** is obvious or “transparent.”
☐ the values of expressions to which **e** refers are obvious or “transparent.”
☐ for all programs **p**, all occurrences of **e** in **p** can be replaced by the result of evaluating **e** without affecting the meaning of **p**.
☐ none of the above

4. (9 pts) **Scala I**

(a) The main programming paradigms of Scala are which of these? (select two)

- ☐ currying ☐ declarative ☐ dysfunctional ☐ functional ☐ object-oriented

(b) What is the result of the following program?

```
val x = 0
def f(y: Int) = y + 1
val result = {
  val x = f(3)
  x * x
} + x
```

- ☐ 0 ☐ 16 ☐ 32 ☐ it does not type check

(c) Why should we care about writing functions that are “tail-recursive?”

- ☐ Recursion should be carried out on the tail, not the head.
☐ Recursion should be carried out on the head, not the tail.
☐ They are “stack-safe”—they help us avoid stack overflows.
☐ They are “disk-safe”—they help us avoid network storage leaks.

5. (6 pts) Let `val X = List(1, 2, 3)`
and `val Y = List(1, 2, 3)`.

(a) To what does the expression `X.map(x => Y.map(y => y - x))` evaluate?

- ☐ `List(0, 0, 0, 0, 0, 0, 0, 0, 0)`
☐ `List(List(0, 0, 0), List(0, 0, 0), List(0, 0, 0))`
☐ `List(0, -1, -2, 1, 0, -1, 2, 1, 0)`
☐ `List(0, 1, 2, -1, 0, 1, -2, -1, 0)`
☐ `List(List(0, 1, 2), List(-1, 0, 1), List(-2, -1, 0))`

(b) To what does the expression `X.flatMap(x => Y.map(y => y - x))` evaluate?

- ☐ `List(0, 0, 0, 0, 0, 0, 0, 0, 0)`
☐ `List(List(0, 0, 0), List(0, 0, 0), List(0, 0, 0))`
☐ `List(List(0, -1, -2), List(1, 0, -1), List(2, 1, 0))`
☐ `List(0, 1, 2, -1, 0, 1, -2, -1, 0)`
☐ `List(List(0, 1, 2), List(-1, 0, 1), List(-2, -1, 0))`

6. (12 pts) **Scala II.** The parts below refer to the function

```
def test(x:Bool, y:Int) = if (x) (y + 2)/y else 0
```

Let CBN = call-by-name

and CBV = call-by-value.

(a) Which strategy evaluates `test(true, 2)` most efficiently (in the fewest steps)?

- ☐ CBN ☐ CBV ☐ CBN and CBV require the same number of steps

(b) Which strategy evaluates `test(true, 1+1)` most efficiently?

- ☐ CBN ☐ CBV ☐ CBN and CBV require the same number of steps

(c) Which strategy evaluates `test(false, 2)` most efficiently?

- ☐ CBV ☐ CBN ☐ CBN and CBV require the same number of steps

(d) Which strategy evaluates `test(false, 1+1)` most efficiently?

- ☐ CBN ☐ CBV ☐ CBN and CBV require the same number of steps

7. (9 pts) **Latency and fault-tolerance.**

(a) *Latency* is degradation in performance due to... (select two)

- ☐ a small number of cores in the central processing unit
☐ slow data transfer across the network or cluster
☐ shuffling data between different nodes in a cluster
☐ stack overflow caused by recursion

(b) Hadoop achieves fault-tolerance by...

- ☐ using lazy evaluation and garbage collection.
☐ writing intermediate computations to disk.
☐ keeping all data immutable and in-memory.
☐ replaying functional transformations over the original (immutable) dataset.

(c) Spark decreases latency while remaining fault-tolerant by... (select three)

- ☐ using ideas from imperative programming.
☐ using ideas from functional programming.
☐ discarding data when it's no longer needed.
☐ keeping all data immutable and in-memory.
☐ replaying functional transformations over the original (immutable) dataset.

8. (12 pts) **Transformations and actions.**

(a) A **transformation** on an RDD... (select two)

- ☐ does not immediately compute a result.
- ☐ immediately computes and returns a result.
- ☐ is lazily evaluated.
- ☐ is eagerly evaluated.

(b) An **action** on an RDD... (select two)

- ☐ does not immediately compute a result.
- ☐ immediately computes and returns a result.
- ☐ is lazily evaluated.
- ☐ is eagerly evaluated.

(c) After performing a series of transformations on an RDD, which of the following methods would ensure that Spark actually carries out the transformations.

- ☐ `mapValues()` ☐ `collect()` ☐ `groupBy()` ☐ none of these

(d) After performing a series of transformations on an RDD, which of the following methods could you use to make sure those transformations are not repeated unnecessarily?

- ☐ `save()` ☐ `persist()` ☐ `collect()` ☐ `parallelize()`

(e) Why does the RDD class have no `foldLeft` method?

- ☐ `foldLeft` is not stack-safe.
- ☐ `foldLeft` is not fault-tolerant.
- ☐ `foldLeft` only works on `PairRDDs`.
- ☐ `foldLeft` is not parallelizable.
- ☐ It's not true; the RDD class *does* have a `foldLeft` method.

(f) Which method of the RDD class has the same effect as `foldLeft` and overcomes limitations of the latter?

- ☐ `aggregate` ☐ `foldRight` ☐ `join` ☐ `leftOuterJoin` ☐ `collect`

9. (12 pts) *Everyday I'm shufflin'.*

(a) What is shuffling?

- ☐ a method for recovering data after hardware failure
- ☐ the method used to ensure a random number generator is unbiased
- ☐ any movement of data
- ☐ moving data from memory to disk, usually caused by insufficient fast memory
- ☐ transferring data between nodes in a cluster, usually in order to complete a computation

(b) What Spark feature or method can be used to reduce or eliminate shuffling?

- ☐ *fault-tolerance*: use higher quality, fault-tolerant hardware
- ☐ *partitioning*: partition an RDD before applying transformations or actions
- ☐ *pre-shuffling*: use a pre-shuffled random number generator
- ☐ *data siloing*: keep the entire RDD on a single node of the cluster
- ☐ *caching*: keep the entire RDD in fast memory on a single node

(c) Which transformations on an RDD of type `RDD[Int]` have we learned about or used in this class?

- ☐ `curry` ☐ `foldLeft` ☐ `groupByKey` ☐ `map` ☐ `mapValues`

(d) Which transformations on a *pair* RDD of type `RDD[(Int, String)]` have we learned about or used in this class? (select three)

- ☐ `curry` ☐ `foldLeft` ☐ `groupByKey` ☐ `map` ☐ `mapValues`

10. (6 pts) **Joins.** Let V and W be types. Suppose `rdd1: RDD[(Int, V)]` and `rdd2: RDD[(Int, W)]` are pair RDDs.

(a) What is the type of `rdd1 join rdd2`.

- ☐ `RDD[(Int, V)]`
- ☐ `RDD[(Int, V \cap W)]`
- ☐ `RDD[(Int, V)]` if `rdd1` is larger than `rdd2`; otherwise, `RDD[(Int, W)]`
- ☐ `RDD[(Int, (V, W))]`
- ☐ `RDD[(Int, (V, Option[W]))]`
- ☐ `RDD[(Int, Option[(V, W)])]`

(b) What is the type of `rdd1 leftOuterJoin rdd2`.

- ☐ `RDD[(Int, V)]`
- ☐ `RDD[(Int, V \cap W)]`
- ☐ `RDD[(Int, V)]` if `rdd1` is larger than `rdd2`; otherwise, `RDD[(Int, W)]`
- ☐ `RDD[(Int, (V, W))]`
- ☐ `RDD[(Int, (V, Option[W]))]`
- ☐ `RDD[(Int, Option[(V, W)])]`

11. (12 pts) **Partitioning.**

Consider a Pair RDD with the following keys: 5, 10, 15, 20, 25, 30, 35, 40, 45. Suppose we partition this RDD into 3 blocks.

- (a) If we use hash partitioning, and if `hashCode()` is the identity (`n.hashCode() == n`), then how many key-value pairs will end up in the first block (block 0) of the partition.

☐ 0 ☐ 1 ☐ 2 ☐ 3 ☐ 4 ☐ 5

- (b) If we use range partitioning with ranges [0, 20], [21, 40], [41, 60], then how many key-value pairs will end up in the first block (block 0) of the partition.

☐ 0 ☐ 1 ☐ 2 ☐ 3 ☐ 4 ☐ 5

- (c) Which strategy results in a more even distribution of the RDD?

☐ (a)'s hash partitioning ☐ (b)'s range partitioning ☐ they're equivalent

- (d) Which of the following transformations, when performed on a partitioned RDD, will result in an RDD with the same partitioning scheme?

☐ `map` ☐ `mapValues` ☐ `foldLeft` ☐ `repartition`