

CS 644: Homework 4

Instructions. Answer the following multiple choice questions by selecting all correct choices.

1. Partitions and Partitioning

- (a) Given a pair RDD (of key-value pairs), when we group values with the same key Spark collects key-value pairs with the same key on the same machine of our cluster.
- ☐ True ☐ False
- (b) By default, Spark uses range partitioning to determine which key-value pair should be sent to which machine.
- ☐ True ☐ False
- (c) Suppose we partition an RDD into a number of blocks. From the following statements, select the two that are true.
- ☐ A single block of the partition may be distributed across multiple machines in the cluster.
 - ☐ A block of the partition is assigned to at most one machine of the cluster.
 - ☐ At least one block of the partition is assigned to every machine in the cluster.
 - ☐ At most one block of the partition is assigned to every machine in the cluster.
 - ☐ More than one block of the partition may be assigned to the same machine in the cluster.

2. Consider a Pair RDD, with keys [8, 23, 39, 40, 97], and suppose we want to partition these data into 4 blocks.

- (a) Using hash partitioning with the identity as `hashCode()` function (`n.hashCode() == n`), check the boxes next to the numbers assigned to the given partition block.
- i. block 0: ☐ 8 ☐ 23 ☐ 39 ☐ 40 ☐ 97 ☐ none
 - ii. block 1: ☐ 8 ☐ 23 ☐ 39 ☐ 40 ☐ 97 ☐ none
 - iii. block 2: ☐ 8 ☐ 23 ☐ 39 ☐ 40 ☐ 97 ☐ none
 - iv. block 3: ☐ 8 ☐ 23 ☐ 39 ☐ 40 ☐ 97 ☐ none
- (b) Using range partitioning with ranges [0, 20], [21, 40], [41, 60], [61, 100], check the boxes next to the numbers assigned to the given partition block.
- i. block 0: ☐ 8 ☐ 23 ☐ 39 ☐ 40 ☐ 97 ☐ none
 - ii. block 1: ☐ 8 ☐ 23 ☐ 39 ☐ 40 ☐ 97 ☐ none
 - iii. block 2: ☐ 8 ☐ 23 ☐ 39 ☐ 40 ☐ 97 ☐ none
 - iv. block 3: ☐ 8 ☐ 23 ☐ 39 ☐ 40 ☐ 97 ☐ none
- (c) Which strategy would result in a more balanced distribution of the data across the partition?
- ☐ hash partitioning ☐ range partitioning

3. (a) Which method can we use to determine whether Spark recognizes that a transformation or action will result in shuffling?
☐ debugDAG ☐ isShuffled ☐ showSchema ☐ showExecutionPlan ☐ toDebugString
- (b) How data is initially partitioned and arranged on the cluster doesn't matter, since Spark will always re-arrange your data to avoid shuffling.
☐ True ☐ False
- (c) `reduceByKey` running on a pre-partitioned RDD will compute values locally, requiring only the final reduced values to be sent from workers to the driver.
☐ True ☐ False
- (d) `join` called on two RDDs that are pre-partitioned with the same partitioner and cached on the same node will cause the join to be computed locally, with no shuffling across the network.
☐ True ☐ False
- (e) Suppose algorithm **A** joins two RDDs and then performs a filter on the result while algorithm **B** performs a filter on the two RDDs and then joins the results. Assume the two algorithms obtain the same result. In general, which algorithm do you expect will cause less data shuffling?
☐ **A** ☐ **B**
4. Answer the following parts by typing in the spaces provided. Select from among the following words or phrases: "at most one," "multiple," "fast," "slow," "some," or "none."
- (a) In a *narrow dependency*, each block of the parent RDD may be used by _____ block(s) of the child RDD.
Narrow dependencies are _____ since they require _____ of the data to be shuffled.
- (b) In a *wide dependency*, each block of the parent RDD may be used by _____ block(s) of the child RDD.
Wide dependencies are _____ since they require _____ of the data to be shuffled.
5. (a) The *query optimizer* of Spark SQL is called
☐ Catalyst ☐ Cobalt ☐ Map Reduce ☐ Platinum ☐ Tungsten
- (b) The *off-heap serializer* of Spark SQL is called
☐ Catalyst ☐ Cobalt ☐ Map Reduce ☐ Platinum ☐ Tungsten

6. (a) Conceptually, DataFrames are RDDs that contain

- ☐ AWS S3 buckets
- ☐ Microsoft Azure blobs
- ☐ Excel spreadsheets
- ☐ Row objects with a known schema
- ☐ Row objects with type information that is checked at compile time

(b) Which of the following can be used to construct a schema identical to the schema that spark would infer if given a collection of objects of type

```
case class Person(name: String, age: Int)?
```

- ☐ Struct(List(Field("name", String), Field("age", Integer)))
- ☐ StructType(List(Field("name", StringType, false),
Field("age", IntegerType, false)))
- ☐ StructType(List(StrucField("name", TypedString)),
List(StructField("age", TypedInteger)))
- ☐ StructType(List(StrucField("name", StringType, true),
StructField("age", IntegerType, true)))
- ☐ Structured(StructuredField("name", String, Boolean) ::
StructuredField("age", Integer, Boolean))

7. (a) Navigate to the Spark API documentation and search for `RelationalGroupedDataset` (the type returned when one calls `groupBy` on a `DataFrame`). Which of the following is **not** a method of the `RelationalGroupedDataset` class?

☐ `agg` ☐ `as` ☐ `avg` ☐ `count` ☐ `min` ☐ `round` ☐ `sum`

(b) Navigate to the Spark API documentation search for `DataFrame`, and notice that none of the results is about the `DataFrame` type itself. This is because

- ☐ `DataFrame` is just an alias for `Dataset[Row]`.
- ☐ `DataFrame` is not a type we really use in Spark or Spark SQL.
- ☐ `DataFrame` is from Spark version 1.0; it is deprecated (no longer supported) in Spark 2.0 or Spark 3.0.
- ☐ `DataFrame` should be spelled `Dataframe`; if you search for `Dataframe` instead, many results appear.

8. (a) `reduceByKey` is a useful method available for RDD's, but is not a method of the `Datasets` class.
- ☐ True ☐ False
- (b) If `reduceByKey` is not available for `Datasets`, which of the following approaches could be used to carry out a Map-reduce operation equivalent to `reduceByKey`?
- ☐ `groupByKey` followed by `mapGroups`
- ☐ `groupByKey` followed by `mapValues` followed by `reduceGroups`
- ☐ `groupByKey` followed by `agg` with a specially constructed `Aggregator` object as argument
- ☐ all of the above
9. (a) If you have unstructured data, you need to fine-tune and manage low-level details of RDD computations, and you have complex data types that cannot be serialized with Encoders, then you should
- ☐ RDDs ☐ DataFrames ☐ Datasets
- (b) If you have structured/semi-structured data and you want the best possible performance, automatically optimized for you, then you should use
- ☐ RDDs ☐ DataFrames ☐ Datasets
- (c) If you have structured/ semi-structured data, you want typesafety, you need to work with functional APIs and you need good performance but it doesn't have to be the best, then you should use
- ☐ RDDs ☐ DataFrames ☐ Datasets